

ADL HW2 Report

editor: M11015Q02 柯元豪

Q1 (Data processing)

Tokenizer

Multiple Choice

在前處理時，會將 question 與 4 個候選的 context 同時做 tokenize，tokenizer 會根據 max_seq_length 將原始句子做斷詞、編碼，其中 tokenizer 加入的 special token 也可能隨著選用的 pretrained model 不同而有所改變，最終會處理成 $[(\text{tokenized_question}, \text{tokenizedcontext}_1), \dots, (\text{tokenized_question}, \text{tokenizedcontext}_4)]$ 的集合

Question Answering

在前處理時，為了避免 question 中有過多空白字元，會先將其移除；接著將處理後的 question 與 context 交給 tokenizer 處理。Tokenizer 會根據 max_seq_length 做 truncation & padding，透過 overflow_to_sample_mapping 能夠將得到的 sequence feature mapping 回去 sequence，透過 offset_mapping 能夠幫助我們 mapping token 所對應到原句子中的 position。

Answer Span

How did you convert the answer span start/end position on characters to position on tokens after BERT tokenization?

Answer Span 階段會將 tokenized 後的 tokenized_examples 逐筆做處理。先取出 input_ids 以此可以得到 CLS_token 所在的 position，也能得到 tokenized sequence length 以利後續尋找 end position。接著先確認此筆 tokenized_example 有沒有給定 answer start 的位置，如果沒有則將 start/end position 都設為 CLS_token 的位置，如果有則將 start position 設為給定的位置，而 end position 則為 start position + answer length，後續再將目前的 start/end position 透過 offset_mapping mapping 回去原句子中的 position。

After your model predicts the probability of answer span start/end position, what rules did you apply to determine the final start/end position?

在 post_processing_function 的步驟，會將模型輸出的 (start_logits, end_logits) 組合，排除所有 start/end position 可能發生的不合理情況，如 end position 比 start position 前面、start position 超過 sequence length ...等，最後排序剩餘組合的分數，取 N 個最高分的 start/end position 儲存在 nbest_predictions.json，最佳的也會直接 mapping 回原句，將答案儲存在 predictions.json。

Q2 (Modeling with BERTs and their variants)

Describe

- configuration of the transformer model

Multiple Choice

```
{
  "_name_or_path": "bert-base-chinese",
  "architectures": [
    "BertForMultipleChoice"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "directionality": "bidi",
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "pooler_fc_size": 768,
  "pooler_num_attention_heads": 12,
  "pooler_num_fc_layers": 3,
  "pooler_size_per_head": 128,
  "pooler_type": "first_token_transform",
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.17.0",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 21128
}
```

Question Answering

```
{
  "_name_or_path": "bert-base-chinese",
  "architectures": [
    "BertForQuestionAnswering"
  ],
  "attention_probs_dropout_prob": 0.1,
```

```
"classifier_dropout": null,  
"directionality": "bidi",  
"hidden_act": "gelu",  
"hidden_dropout_prob": 0.1,  
"hidden_size": 768,  
"initializer_range": 0.02,  
"intermediate_size": 3072,  
"layer_norm_eps": 1e-12,  
"max_position_embeddings": 512,  
"model_type": "bert",  
"num_attention_heads": 12,  
"num_hidden_layers": 12,  
"pad_token_id": 0,  
"pooler_fc_size": 768,  
"pooler_num_attention_heads": 12,  
"pooler_num_fc_layers": 3,  
"pooler_size_per_head": 128,  
"pooler_type": "first_token_transform",  
"position_embedding_type": "absolute",  
"torch_dtype": "float32",  
"transformers_version": "4.17.0",  
"type_vocab_size": 2,  
"use_cache": true,  
"vocab_size": 21128  
}
```

- performance

Multiple Choice

```
# Validation Set  
Accuracy: 0.9937
```

Question Answering

```
# Test Set  
EM: 0.7749 (Kaggle)  
  
# Validation Set  
EM: 0.9351
```

- loss function

Multiple Choice

CrossEntropyLoss

Question Answering

CrossEntropyLoss

- The following hyperparameters were used during training:

Multiple Choice

```
learning_rate: 3e-05
train_batch_size: 1
eval_batch_size: 8
seed: 42
gradient_accumulation_steps: 2
total_train_batch_size: 2
optimizer: Adam with betas=(0.9,0.999) and epsilon=1e-08
lr_scheduler_type: linear
num_epochs: 2.0
```

Question Answering

```
learning_rate: 3e-05
train_batch_size: 4
eval_batch_size: 8
seed: 42
gradient_accumulation_steps: 2
total_train_batch_size: 8
optimizer: Adam with betas=(0.9,0.999) and epsilon=1e-08
lr_scheduler_type: linear
num_epochs: 2.0
```

Try another type of pretrained model and describe

Question Answering (Use the same Multiple Choice model)

- configuration of the transformer model

```
{
  "_name_or_path": "hfl/chinese-roberta-wwm-ext",
  "architectures": [
    "BertForQuestionAnswering"
  ],
  "attention_probs_dropout_prob": 0.1,
  "bos_token_id": 0,
  "classifier_dropout": null,
  "directionality": "bidi",
  "eos_token_id": 2,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "output_past": true,
  "pad_token_id": 0,
  "pooler_fc_size": 768,
  "pooler_num_attention_heads": 12,
  "pooler_num_fc_layers": 3,
  "pooler_size_per_head": 128,
  "pooler_type": "first_token_transform",
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.17.0",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 21128
}
```

- performance

```
# Test Set
EM: 0.7966 (Kaggle)

# Validation Set
EM: 0.9461
```

- difference

architecture

	bert-base-chinese	chinese-roberta-wwm-ext
Masking	WordPiece	WWM ^[1]
Type	base	base
Data Source	wiki	wiki+ext ^[2]
Training Tokens #	0.4B	5.4B
Device	TPU Pod v2	TPU v3
Optimizer	AdamW	AdamW
Vocabulary	21,128	~bert-base-chinese
Init Checkpoint	Random Init	~bert-base-chinese

[1] WWM: Whole Word Masking [2] ext: extended data


RoBERTa 與 BERT 不同在於，RoBERTa 取消了 BERT 原有的 Next Sentence Prediction(NSP) 任務，並將原本靜態 Masking 的方法改成動態的 Masking。

pretraining loss

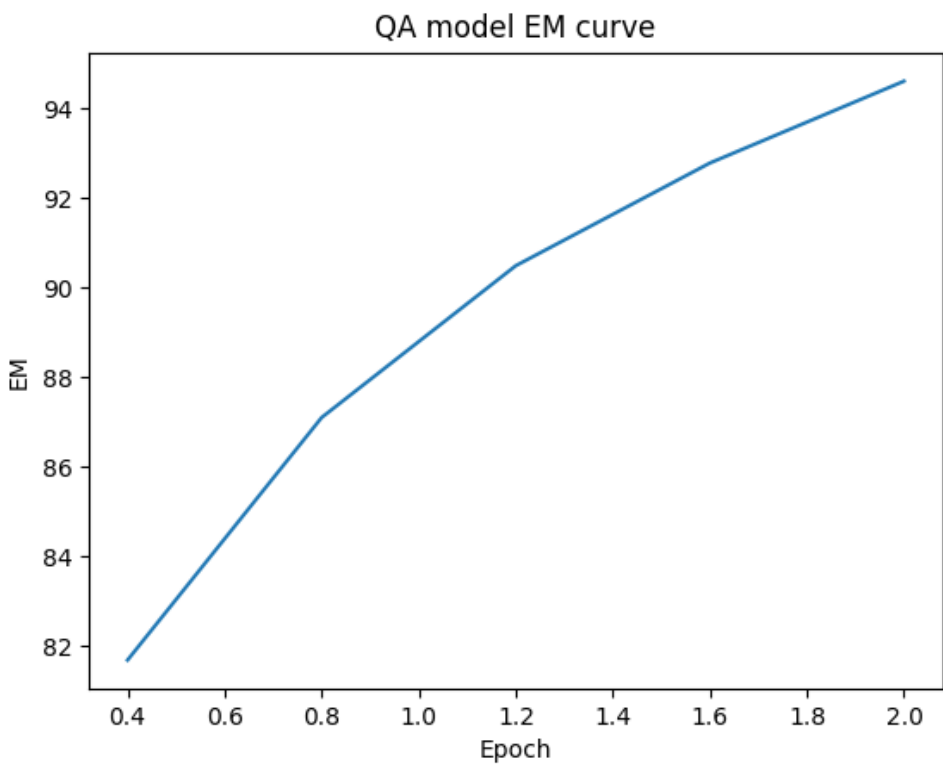
```
# bert-base-chinese
"train_loss": 0.7986815468019802
# chinese-roberta-wwm-ext
"train_loss": 0.7193350367217894
```

Q3 (Curves)

Learning curve of loss



Learning curve of EM



Q4 (Pretrained vs Not Pretrained)

Configuration of the model and how do you train this model

- config

```
{
  "_name_or_path": "tmp/qa_chinese_roberta_wwm_ext/",
  "architectures": [
    "BertForQuestionAnswering"
  ],
  "attention_probs_dropout_prob": 0.1,
  "bos_token_id": 0,
  "classifier_dropout": null,
  "directionality": "bidi",
  "eos_token_id": 2,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "output_past": true,
  "pad_token_id": 0,
  "pooler_fc_size": 768,
  "pooler_num_attention_heads": 12,
  "pooler_num_fc_layers": 3,
  "pooler_size_per_head": 128,
  "pooler_type": "first_token_transform",
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.17.0",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 21128
}
```


- how to train

將

```
model = AutoModelForQuestionAnswering.from_pretrained()
```

替換成：

```
config = AutoConfig.from_pretrained(model_config_path)
```

```
model = AutoModelForQuestionAnswering.from_config(config)
```

從相同模型設定的資料夾讀取 config, AutoModel 中的 from_pretrained 會 load weight,

而 from_config 則不會，其餘的 training script 不變，即可 training from scratch

Performance of this model v.s. BERT

```
# Evaluate on valid set
```

```
(Without pretrained weight)
```

```
{  
  "eval_exact_match": 6.945829179129279,  
  "eval_f1": 6.945829179129279,  
}
```

```
(With pretrained weight)
```

```
{  
  "eval_exact_match": 94.6161515453639,  
  "eval_f1": 94.6161515453639,  
}
```

Q5 (Bonus: HW1 with BERTs)

model

intent

```
{
  "_name_or_path": "bert-base-cased",
  "architectures": [
    "BertForSequenceClassification"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "id2label": {
    "0": "accept_reservations",
    "1": "account_blocked",
    "2": "alarm",
    .
    {太冗長，故省略}
    .
    "147": "who_do_you_work_for",
    "148": "who_made_you",
    "149": "yes"
  },
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "label2id": {
    "accept_reservations": 0,
    "account_blocked": 1,
    "alarm": 2,
    .
    {太冗長，故省略}
    .
    "who_do_you_work_for": 147,
    "who_made_you": 148,
    "yes": 149
  },
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "problem_type": "single_label_classification",
  "torch_dtype": "float32",
  "transformers_version": "4.17.0",
```

```
"type_vocab_size": 2,  
"use_cache": true,  
"vocab_size": 28996  
}
```

slot tagging

```
{  
  "_name_or_path": "bert-base-cased",  
  "architectures": [  
    "BertForTokenClassification"  
  ],  
  "attention_probs_dropout_prob": 0.1,  
  "classifier_dropout": null,  
  "finetuning_task": "ner",  
  "gradient_checkpointing": false,  
  "hidden_act": "gelu",  
  "hidden_dropout_prob": 0.1,  
  "hidden_size": 768,  
  "id2label": {  
    "0": "B-date",  
    "1": "B-first_name",  
    "2": "B-last_name",  
    "3": "B-people",  
    "4": "B-time",  
    "5": "I-date",  
    "6": "I-people",  
    "7": "I-time",  
    "8": "O"  
  },  
  "initializer_range": 0.02,  
  "intermediate_size": 3072,  
  "label2id": {  
    "B-date": 0,  
    "B-first_name": 1,  
    "B-last_name": 2,  
    "B-people": 3,  
    "B-time": 4,  
    "I-date": 5,  
    "I-people": 6,  
    "I-time": 7,  
    "O": 8  
  },  
  "layer_norm_eps": 1e-12,  
  "max_position_embeddings": 512,  
  "model_type": "bert",  
  "num_attention_heads": 12,  
  "num_hidden_layers": 12,  
  "pad_token_id": 0,  
  "position_embedding_type": "absolute",  
}
```

```

"torch_dtype": "float32",
"transformers_version": "4.17.0",
"type_vocab_size": 2,
"use_cache": true,
"vocab_size": 28996
}

```

performance

intent classification

Submission and Description	Private Score	Public Score
intent_submission.csv just now by YuanHao bert-base	0.96044	0.95955

slot tagging

Submission and Description	Private Score	Public Score
slot_submission.csv 3 hours ago by YuanHao bert-base	0.83869	0.83699

Loss function

intent classification & slot tagging

CrossEntropyLoss

The optimization algorithm (e.g. Adam), learning rate and batch size.

intent classification & slot tagging

```

learning_rate: 5e-05
train_batch_size: 8
eval_batch_size: 8
seed: 42
optimizer: Adam with betas=(0.9,0.999) and epsilon=1e-08
lr_scheduler_type: linear
num_epochs: 3.0

```